Fundamentals of Reinforcement Learning

Luka Meladze
National University of Singapore
CP2107 Project: **LLM Post-training with Reinforcement Learning**

Contents

Abstract

This paper provides an oreview of Reinforcement Learning (RL) concepts, bridging fundamental principles with their modern application in post-training Large Language Models (LLMs). It establishes theoretical groundwork, defines the Markov Decision Process (MDP) and its components, and distinguishes RL from Supervised and Imitation Learning, highlighting RL's unique ability to learn through environmental interaction rather than from static, labeled datasets. The report also details the mathematical formulation of the RL objective across finite and infinite horizons, focusing on derivation of the policy gradient theorem and policy gradient methods, discussing essential variance reduction techniques like causality-based reward-to-go and baselines, and presenting a simplified algorithmic implementation.

The paper then transitions to the practical application of these concepts in the LLM pipeline, examines the evolution of an LLM from a base model trained on next-token prediction to a helpful assistant through Supervised Fine-Tuning (SFT), and analyzes the inherent limitations of SFT language models, including their token-based architecture and finite per-token computational power, which can lead to failures in seemingly simple tasks. Subsequently, we explore how RL, particularly through direct policy optimization in verifiable domains (e.g., mathematics) and Reinforcement Learning from Human Feedback (RLHF) in subjective domains, overcomes these limitations to move beyond mere imitation and discover effective reasoning strategies.

1 Reinforcement Learning

1.1 Markov Decision Process (MDP)

A **Markov Decision Process (MDP)** is a framework used to describe sequential decision-making problems, where a decision-maker (agent) interacts with an environment over time. At each time step t, the agent observes an observation of the state of the world s_t (possibly partial o_t), takes an action a_t, and then receives a reward r_t while making a transition to the new state s_{t+1}. This cycle continues, and the agent's ultimate goal is to find a **policy** (a strategy for choosing actions) that maximises the total expected reward.

The key idea and assumption of MDP is rooted in the **Markov property**, which states that the next state depends **only** on the current state and action, not on the history of previous states or actions. That doesn't mean that the future is perfectly predictable; there might still be randomness, but knowing the past doesn't help at all to resolve that randomness.

An MDP is defined by a **5-tuple**: $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$.

- **States ($\mathcal{S}$)**: The set of all possible states the environment of an agent can be in. Each state $s \in \mathcal{S}$ is a **complete** description of the environment, while an observation o is a partial description of a state. A state s must satisfy the **Markov Property**: If the current state s_t is given, the future state s_{t+1} and future rewards are independent of past states $(s_1, \ldots, s_{t-1})$ and actions $(a_1, \ldots, a_{t-1})$. Mathematically, $P(s_{t+1}|s_t, a_t, \ldots, s_1, a_1) = P(s_{t+1}|s_t, a_t)$.

- **Actions ($\mathcal{A}$)**: The set of all possible & valid actions (moves or decisions) that the agent can make in any state is often called the action space. Some environments have discrete action spaces, while others have continuous action spaces where actions are real-valued vectors.

- **Transition Probability Operator/Function(T)**: It is a **tensor** which specifies the probability of transitioning from the current state s to the next state s' after taking action a. $T_{i,j,k} = p(s_{t+1} = i|s_t = j, a_t = k)$ also denoted as $P(s'|s, a)$. The environ-

73 ment's response to an action can be **deterministic** (always leads to the same next
74 state) or **stochastic** (leads to different next states).

75 • **Reward Function (R)**: The reward function R is very important in reinforcement
76 learning, which is a mapping from $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$. It depends on the current state
77 of the world and the action just taken. It quantifies the desirability of an action
78 or state, typically denoted as $R(s)$ or $R(s, a)$. The agent's ultimate goal in RL is to
79 maximise **not** *the immediate* but the *cumulative* reward over time (over a trajectory).

80 • **Discount Factor (γ)**: For tasks that continue indefinitely, the cumulative sum of
81 rewards could be infinite. We need ways to ensure the objective remains finite.
82 Discount Factor is a value between 0 and 1 ($\gamma \in [0, 1)$ typically 0.9 to 0.99) which
83 weights rewards received sooner more heavily than rewards received later to guar-
84 antee that the infinite sum converges to a finite value.

1.2 How RL Differs from Supervised / Imitation Learning

86 Reinforcement Learning (RL) and Supervised Learning (SL), including Imitation Learning
87 (IL), learn in fundamentally different ways and are distinguishable by their data sources,
88 goals, inherent assumptions, and feedback mechanisms. The Supervised Learning works
89 fine under the assumption that each data example is Independent and Identically Dis-
90 tributed. **Imitation Learning (Behavioural Cloning)** is a form of Supervised Learning
91 applied to sequential decision-making. It trains a policy π_θ to mimic an expert's actions.
92 However, IL inherits SL's limitations, particularly the distributional shift problem, meaning
93 it cannot guarantee optimal performance outside the expert's observed trajectory space and
94 cannot discover behaviours superior to the expert's. RL directly addresses these limitations
95 by learning through interaction and optimising for long-term outcomes.

96 • In supervised learning, including imitation learning, the agent learns from a fixed
97 dataset (**Independent and Identically Distributed data**) of input-output pairs
98 $\mathcal{D} = \{(X_i, Y_i)\}_{i=1}^{N}$. The model learns from these static examples and receives direct,
99 immediate, and explicit "ground truth" labels (correct answers) for each input.

101 In contrast, reinforcement learning does not assume access to an expert or labelled
102 actions. Instead, the agent sequentially interacts with the environment directly. It
103 observes the state, selects actions, and receives rewards based on those actions. The
104 agent's goal is not to imitate an expert but to maximise the total cumulative reward
105 it receives over time.

106 • **Learning Objective**: In supervised learning, the goal is to generalise patterns from
107 given data and copy the behaviour or labels. i.e. learn mapping $f_\theta : \mathcal{X} \to \mathcal{Y}$ that
108 accurately predicts outputs for unseen inputs and minimises prediction error.

109 In Reinforcement Learning, the goal is to learn a **policy** $\pi_\theta : \mathcal{S} \to \mathcal{A}$ that maximises
110 the **total future reward** over a sequence of interactions. It's about figuring out what
111 the agent should do to reach a goal, potentially discovering solutions better than
112 any human expert could show (like AlphaGo's surprising "Move 37").

113 It maximises $J(\theta) = \mathbb{E}_{\tau \sim P_\pi(\tau)} \left[\sum_{t=1}^{T} \gamma^t R(s_t, a_t) \right]$.

1.3 Real-World Case: Autonomous Drone trained by RL

115 Let's give an example and describe a real-world case of an autonomous drone that navigates
116 through a forest or different obstacles to reach a destination while avoiding any barriers.
117 The AI-powered drone's goal is to learn a policy that maximises its total reward, meaning it
118 learns to reach its destination safely while minimising crashes and being efficient.

119 • **Agent**: The drone's control system (the AI brain that makes decisions about flying).

120 • **The state** s_t at time t includes all relevant information about the environment, such
121 as the drone's position, speed, orientation, battery level, and a camera image (the
122 pixels showing trees, clear paths, etc.) of its surroundings.

3

- **Actions ($\mathcal{A}$)**: These are the commands the drone can give to its motors to change its flight. The action a_t could be continuous control commands like pitch and roll or **discrete** commands like "turn left," "turn right," or "go straight."

- **Reward Function (R)**: How the drone gets feedback on its performance to learn and improve:

    - **Positive Reward**: For moving closer to the destination, or for each second it stays flying without crashing.
    - **Negative Reward (Penalty)**:
        * A large negative reward for colliding with a tree or any other obstacle.
        * A small negative reward for using a lot of battery power.

- **Transition Probability Function ($P(s_{t+1}|s_t, a_t)$)**: This describes how the drone moves through the forest based on the current state when an action is applied, including sensor noise and wind disturbances. This is typically **stochastic** due to external factors like wind.

- **Policy (π_θ)**: This is the learned strategy or "brain" of the drone. It tells the drone what action to take in the current state (e.g., "When it sees a particular state/image of obstacles ahead, tilt hard right!").

- **Trajectory (τ)**: A sequence of states and actions along the way. $\tau = (s_1, a_1, s_2, a_2 \ldots)$.

- **Objective**: The objective is to learn policy parameters θ^* that maximise the **expected cumulative reward** over extended periods and to train the drone to learn this policy so that it can reach the destination safely and efficiently.

1.4 The Objective of Reinforcement Learning in Math

The fundamental objective in Reinforcement Learning is to find an **optimal policy (π^*)** that is defined by parameters (θ^*) and maximises the **expected cumulative reward** (also known as "return") over time. This expectation accounts for the stochasticity of both the environment and the agent's policy.

1.4.1 Finite Horizon (Total Reward)

For tasks that have a fixed, finite number of time steps T, the objective is to maximise the expected sum of rewards the agent collects until the end of completing the task.

The **return** for a single trajectory $\tau = (s_1, a_1, \ldots, s_T, a_T, s_{T+1})$ is:

$$G_0 = \sum_{t=1}^{T} R(s_t, a_t)$$

The **objective function**, which we aim to maximise by optimising the policy's parameters θ, is the **expected return**:

$$J(\theta) = \mathbb{E}_{\tau \sim P_\pi(\tau)} \left[\sum_{t=1}^{T} R(s_t, a_t) \right]$$

Where:

- $P_\pi(\tau)$: The probability of a specific trajectory τ occurring under policy π_θ:

$$P_\pi(\tau) = P(s_1) \prod_{t=1}^{T} \pi_\theta(a_t|s_t) P(s_{t+1}|s_t, a_t)$$

- $R(s_t, a_t)$: The immediate reward received at time t in state s_t taking an action a_t.

158 An equivalent formulation for finding the optimal policy parameters θ^* in the finite horizon
159 case is to maximise the sum of expected rewards at each time step. Here, $p_\theta(s_t, a_t)$ represents
160 the probability of visiting state s_t and taking action a_t at time t under policy π_θ.

$$\theta^* = \arg\max_\theta \sum_{t=1}^{T} \mathbb{E}_{(s_t, a_t) \sim p_\theta(s_t, a_t)} [R(s_t, a_t)]$$

1.4.2 Infinite Horizon (Average Reward)

162 For tasks where the agent operates continuously without a natural ending point and without
163 a discount factor, the objective is to maximise the long-term **average** reward per time step.

164 The **average reward objective** is defined as:

$$J_{avg}(\theta) = \lim_{T \to \infty} \frac{1}{T} \mathbb{E}_{\tau \sim P_\pi(\tau)} \left[\sum_{t=1}^{T} R(s_t, a_t) \right]$$

165 When the system converges to a stationary distribution $p_\theta(s, a)$, this objective simplifies to
166 the expected immediate reward under that distribution:

$$J_{avg}(\theta) = \mathbb{E}_{(s,a) \sim p_\theta(s,a)} [R(s, a)]$$

167 The optimal policy parameters θ^* for the average reward case are therefore found by
168 maximising this expected immediate reward under the stationary distribution.

$$\theta^* = \arg\max_\theta \mathbb{E}_{(s,a) \sim p_\theta(s,a)} [r(\mathbf{s}, \mathbf{a})]$$

1.4.3 Infinite Horizon (Discounted Total Reward)

170 For continuous tasks without a natural ending point, we also use a **discounted sum of**
171 **rewards** to ensure the total return remains finite and converges.

172 The **discounted return** for a single trajectory $\tau = (s_1, a_1, s_2, a_2, \dots)$ is:

$$G_0 = \sum_{t=1}^{\infty} \gamma^t R(s_t, a_t)$$

173 The **objective function** to maximise is the **expected discounted return**:

$$J(\theta) = \mathbb{E}_{\tau \sim P_\pi(\tau)} \left[\sum_{t=1}^{\infty} \gamma^t R(s_t, a_t) \right]$$

174 Where:

175 • γ: The **discount factor**, a value in $[0, 1)$ that exponentially weighs future rewards
176   less than immediate ones.

177 The ultimate goal in Reinforcement Learning is to find the policy parameters θ^* that max-
178 imise this objective function:

$$\theta^* = \arg\max_\theta J(\theta)$$

179 This means we are searching for the best possible strategy that, on average, leads to the
180 highest accumulation of rewards over the long run, considering both the agent's actions
181 and the environment's stochastic (random) nature.

1.5 Policy Gradients

1.5.1 Deriving the Basic Policy Gradient

To achieve the main goal in Reinforcement Learning, i.e. to find the best settings (parameters θ) for a policy π_θ so that the average total reward $J(\theta) = \mathbb{E}_{\tau \sim P_\pi(\tau)} \left[\sum_{t=1}^{T} R(s_t, a_t) \right]$ is as high as possible, policy gradient methods are used that directly optimize a parameterized policy. The underlying principle of policy gradients involves iteratively adjusting policy parameters to increase the probabilities of actions that lead to higher returns while simultaneously decreasing the probabilities of actions that result in lower returns. This iterative refinement process inherently embodies a "trial and error" (Levine, 2021) learning paradigm, where the agent continuously adapts its policy based on observed outcomes. The optimisation of $J(\theta)$ is achieved through an iterative process known as **gradient ascent**, that calculates "steepest way up" given by the **gradient** of the objective function, $\nabla_\theta J(\theta)$. Once this gradient is known, policy parameters θ are adjusted in that direction:

$$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$$

Here, α (alpha) is a small step size, called the **learning rate**.

The challenge is to calculate $\nabla_\theta J(\theta)$ because the trajectory τ depends on the rules of the environment, such as the <u>unknown</u> transition probabilities $P(s_{t+1}|s_t, a_t)$ or initial state $P(s_1)$. We can only interact with the world to collect samples. Mathematically, $J(\theta)$ can be expressed as an average over all possible trajectories τ:

$$J(\theta) = E_{\tau \sim p_\theta(\tau)}[r(\tau)] = \int p_\theta(\tau) r(\tau) d\tau$$

Next, taking the gradient with respect to θ, since the gradient operation is linear, it can be moved inside the integral:

$$\nabla_\theta J(\theta) = \nabla_\theta \int p_\theta(\tau) r(\tau) d\tau = \int \nabla_\theta p_\theta(\tau) r(\tau) d\tau$$

Now, a very useful trick called the "convenient identity" (log-derivative trick) is employed. This trick states that for any probability distribution $P(\tau)$:

$$\nabla_\theta P(\tau) = P(\tau) \nabla_\theta \log P(\tau)$$

Applying this convenient identity is useful because it essentially converts a gradient of a probability into a probability times the gradient of its logarithm. The equation becomes:

$$\nabla_\theta J(\theta) = \int p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) r(\tau) d\tau = E_{\tau \sim p_\theta(\tau)}[\nabla_\theta \log p_\theta(\tau) r(\tau)]$$

Now, let's determine $\nabla_\theta \log p_\theta(\tau)$. The probability of a trajectory is defined as:

$p_\theta(\tau) = p(s_1) \prod_{t=1}^{T} \pi_\theta(a_t|s_t) p(s_{t+1}|s_t, a_t)$

Taking the logarithm of both sides (which turns products into sums) gives the following equation:

$$\log p_\theta(\tau) = \log p(s_1) + \sum_{t=1}^{T} \log \pi_\theta(a_t|s_t) + \sum_{t=1}^{T} \log p(s_{t+1}|s_t, a_t)$$

A crucial simplification occurs when taking the gradient ∇_θ with respect to θ, since the terms $\log p(s_1)$ (initial state distribution) and $\log p(s_{t+1}|s_t, a_t)$ (environment dynamics & transition probabilities) vanish, as they are independent of the policy parameters θ. This is

because the environment's rules do not depend on the policy parameters θ. So, only the policy term remains:

$$\nabla_\theta \log p_\theta(\tau) = \sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_t|s_t)$$

This means the gradient computation now only requires differentiating the known, parameterised policy $pi_\theta(a_t|s_t)$, not the unknown environment. Substituting this back into the policy gradient equation yields the fundamental **basic policy gradient formula** (the REINFORCE Algorithm). This seminal algorithm was formally introduced by (Williams, 1992), providing a foundational method for estimating the policy gradient through a sample mean derived from collected trajectories.

$$\nabla_\theta J(\theta) = E_{\tau \sim p_\theta(\tau)} \left[\left(\sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_t|s_t) \right) \left(\sum_{t=1}^{T} r(s_t, a_t) \right) \right]$$

In practice, since this expectation cannot be calculated perfectly, it is estimated **(unbiased)** by running the policy π_θ in the environment N times and collecting many trajectories. Then, for each sampled trajectory, the calculated terms are averaged to get an estimate of the gradient. So, the **MONTE Carlo approximation formula for policy gradient (MIT-6.7920, 2024)**:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} \left[\left(\sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_{i,t}|s_{i,t}) \right) \left(\sum_{t=1}^{T} r(s_{i,t}, a_{i,t}) \right) \right]$$

1.5.2 Adding a Baseline for the Policy Gradient and Why It is Needed

The basic policy gradient, while mathematically correct, has a significant practical issue: it can be very **noisy** and have **high variance** (Afachao, 2023). This means that if a small number of trajectories are collected, the estimate of the gradient can jump around a lot, making it hard for policy parameters to learn & improve consistently. Additionally, if all the rewards are positive, even for "bad" trajectories, the basic policy gradient method will try to increase the probability of all sampled actions, simply because their reward is positive. This goes against the intuition that the goal is to increase probabilities for "better than average" actions and decrease them for "worse than average" ones.

To avoid this problem and make learning more stable and efficient, <u>two</u> key tricks, reward-to-go and baselines (Afachao, 2023) help reduce the **high variance** of policy gradients:

Exploiting Causality - "Reward-to-Go" ($\hat{Q}_t$)   An action taken at time t cannot affect rewards that were received in the past (before time t) (Schulman, 2016). This is the principle of **causality**. So, when evaluating how good an action a_t was, only the rewards that came **after** that action, from time t until the end of the trajectory, need to be considered. This sum of future rewards is called the **reward-to-go**:

$$\hat{Q}_t = \sum_{t'=t}^{T} r(s_{t'}, a_{t'})$$

Using this reward-to-go term instead of the total reward from the entire trajectory reduces noise by discarding irrelevant past rewards. The policy gradient now looks like:

$$\nabla_\theta J(\theta) \approx E_{\tau \sim p_\theta(\tau)} \left[\sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_t|s_t) \hat{Q}_t \right]$$

Baselines $b(s_t)$ To further reduce noise and make the learning more intuitive, a value called a **baseline**, $b(s_t)$, is subtracted from the reward-to-go term. This baseline typically represents the average or expected reward-to-go from a given state s_t. The use of such a baseline is formally justified by the EGLP (Expected Gradient of Log Policy) Lemma, which states that adding a baseline $b(s_t)$ that is independent of the action a_t does not change the expectation of the policy gradient, while significantly reducing its variance (OpenAI, Achiam, 2018).

By using a baseline, the actual reward received ($\hat{Q}_t$) is essentially compared to what was expected from that state ($b(s_t)$).

- If ($\hat{Q}_t - b(s_t)$) is positive, it means the action taken led to a **better-than-expected** outcome. So, the policy will be adjusted to make that action more likely.

- If ($\hat{Q}_t - b(s_t)$) is negative, it means the action led to a **worse-than-expected** outcome. So, the policy will be adjusted to make that action less likely.

This "centering" provides a more stable learning signal, leading to more reliable policy improvement. So, the formula (MIT-6.7920, 2024) with reward-to-go and a baseline is:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} \sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_{i,t}|s_{i,t}) \left(\hat{Q}_{i,t} - b \right)$$

1.5.3   On-policy algorithm for the Policy Gradient

This algorithm for the policy gradient outlines the steps for how an agent learns using policy gradients with a basic baseline. It focuses on the "trial and error" learning process that policy gradients formalise, allowing an agent to learn complex behaviours by continually refining its strategy based on the outcomes it experiences.

Algorithm: Policy Gradient Method (Simplified version of REINFORCE algorithm (Achiam, 2018))

1. **Input**:

   - Initial policy parameters θ_0 (the starting settings for the "brain")
   - Learning rate α (how big of a step to take when adjusting settings)

2. **For** $k = 0, 1, 2, \ldots$ (repeat for many learning cycles) **do**:

   (a) **Collect Trajectories (Run the Policy)**:
      - Use the current policy $\pi(\theta_k)$ to interact with the environment.
      - Collect a set of trajectories $\mathcal{D}_k = \{\tau_i\}_{i=1}^{N}$ (many paths the robot takes). Each τ_i includes a sequence of states, actions, and rewards: $(s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \ldots, s_{i,T}, a_{i,T}, r_{i,T})$.

   (b) **Compute Rewards-to-Go ($\hat{Q}_t$)**:
      - For each time step t in every collected trajectory $\tau_i \in \mathcal{D}_k$, calculate the reward-to-go $\hat{Q}_{i,t}$:

$$\hat{Q}_{i,t} = \sum_{t'=t}^{T} r(s_{i,t'}, a_{i,t'})$$

   (c) **Compute Simple Baseline (b)**:
      - Calculate the average of all $\hat{Q}_{i,t}$ values across all time steps in all trajectories in the current batch $\mathcal{D}_k$. This average will be our baseline b:

$$b = \frac{1}{\sum_{\tau \in \mathcal{D}_k} |\tau|} \sum_{\tau \in \mathcal{D}_k} \sum_{t=1}^{|\tau|} \hat{Q}_{i,t}$$

      (where $|\tau|$ is the length of trajectory τ)

8

(d) **Estimate Policy Gradient ($\hat{g}_k$)**:
  • Calculate the estimated policy gradient $\hat{g}_k$ using the collected data, reward-to-go values, and the baseline:

$$\hat{g}_k = \frac{1}{N} \sum_{i=1}^{N} \sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(a_{i,t}|s_{i,t})(\hat{Q}_{i,t} - b)$$

(e) **Update Policy Parameters**:
  • Adjust the policy parameters θ by taking a step in the direction of the estimated gradient:

$$\theta_{k+1} = \theta_k + \alpha \hat{g}_k$$

2 Overview and Deep Dives into LLMs like ChatGPT

This section provides an overview of the principles, architecture, and training methodologies behind Large Language Models (LLMs) like ChatGPT and the role of Reinforcement Learning (RL) in improving their outputs. The primary source of the concepts and interesting LLM insights mentioned below is the YouTube Video by Andrej Karpathy (Karpathy, 2025). The report outlines the entire LLM pipeline, from the initial pre-training on vast internet datasets to the post-training stages of Supervised Fine-Tuning (SFT) and RL. The SFT models have inherent computational limitations, and we will explore the specific reasons "why?" they excel at certain (even very hard for humans) tasks while failing at much simpler ones (for humans) like counting or spelling. Finally, we will see how RL is used to unlock reasoning capabilities, moving these models from simple text predictors to more sophisticated problem-solving agents.

2.1 Base-Model

Before a model can learn to "think" for itself or answer questions to become an "assistant", it must first process a huge amount of diverse data in different languages and knowledge sources. This training is a two-stage process: First, training a base model, and then turning this base model into a sophisticated assistant.

- 1.1. The first step in pre-training is data curation, usually from the public internet. The raw data is noisy and requires filtering like URL filtering, Text Extraction, Language Filtering, and PII Removal. After this extensive cleaning, the vast internet is condensed into a more manageable & accurate dataset. For instance, the FineWeb dataset contains approximately 44 terabytes of text (15 trillion tokens), which can then be used as the "textbook" from which the model will learn and be trained. (Penedo et al., 2024)

- 1.2 Tokenisation: Neural networks do not process raw text or bytes. They process sequences of numbers, i.e. tokens (OpenAI, 2025). Tokenisation is the crucial step of converting a stream of text into a sequence of integers (token = Unique ID). Instead of using individual characters (which would create very long sequences) or whole words (which would create an unlimited vocabulary), LLMs use an algorithm like **Byte-Pair Encoding (BPE)** (Hugging Face, 2025) to find commonly occurring chunks of text and assign them a unique token ID. This results in a fixed vocabulary. Tokenisation is important for understanding LLM limitations. The model does not see individual letters; it sees a sequence of text chunks (tokens).

- 1.3 This tokenised data from a diverse collection of texts is then used to train a massive neural network (a Transformer). The initial base model's **objective** is simple: to predict the next token in a sequence. To train an initial base model:
  - The model is given a chunk of text (a "context window"), for example, "The capital of **Georgia** is". (Of course, using tokens like [464, 3139, 286, 7859, 318])
  - It then predicts the most probable next token (in this case, likely the tokens for " Tbilisi"). The model makes its prediction, which is initially random.

- It compares its prediction to the actual next token in the training data and then, using a process called backpropagation, it calculates its error and slightly adjusts its internal parameters (billions of "weights") to make its prediction for that specific example slightly better next time.

By doing this trillions of times in parallel (using tens of thousands of GPUs running for months), the model develops a base model—a massive neural network (like Meta's Llama 3 or GPT-4's base) whose parameters have been tuned to be a "internet document simulator", but is not yet a helpful assistant. If you prompt it, it will simply continue your text as if completing a webpage it found online.

2.2 From Base-Model to Assistant - Supervised Fine-Tuning (SFT)

The base model can only statistically complete texts but is not capable of having conversations and answering questions. To make it a helpful assistant, we need to teach it a new behaviour: how to have a conversation. This is done by further using new tokens specifically for signalling the start & end of conversations, and training models on a much smaller, high-quality dataset of example <u>conversations</u>. In these conversations, human labellers & experts (following detailed instructions & templates) have written "ideal" responses to a wide variety of prompts. It's not learning new factual knowledge (that came from pre-training), but rather a new style of interaction. So, when we ask LLM simple models like GPT-2 or GPT-4o a question, they essentially give back the style of a statistical simulation of what an average human labeller, following OpenAI's guidelines, would have written.

2.3 Errors & Limitations: Why LLMs Can't Spell or Count Characters?

2.3.1 Models are not good at "mental" computations

To generate each token, the model performs a single forward pass through its neural network. This means it takes the current sequence of tokens as input and, after a fixed number of computational layers (e.g., 100 layers), outputs the probabilities for the next token. The critical implication is that the model has a limited computational budget for "thinking" about each token it produces. It cannot loop, ponder indefinitely, or perform a complex multi-step calculation in its "head" before producing a single token. All reasoning & computation must be spread across the multiple steps of generating next tokens.

2.3.2 Why LLMs Can't Spell or Count Characters?

The two characteristics of LLMs: 1. Token-based view & 2. finite computation per token, cause them to have some holes in performing basic tasks, even though they are usually very good at much more advanced problem-solving. As models don't see letters, and they see tokens, the words might be seen as the tokens of text chunks that the model must first recall (from its training), then has to recall the "length" or symbols of these tokens and e.g. add them up. This is a multi-step task that is not native to its architecture, especially in one pass. It's doing "mental arithmetic" on the composition of tokens and often fails. Example prompts where it fails due to tokenisation include **counting** the number of dots, spelling, etc. This is why asking it to use a tool like the Python code interpreter often works: the model can better copy-paste the string (easy task for LLM) and let Python do the actual counting.

2.4 Reinforcement Learning (RL) for LLMs

While SFT creates a competent assistant, it's limited by imitation and its training data. The model only stochastically imitates how a human would write a solution based on the internet data. It doesn't truly "solve" the problem itself. RL addresses this by allowing the model to practice. Instead of just imitating a "perfect" (expert-based, training data-driven) solution, the model generates thousands of its own potential solutions and is rewarded

Figure 1: GPT-4o

for the ones that reach the "correct" or "better" answers. This practice of models without human supervision is important because the optimal reasoning path for an LLM (a sequence of tokens) is often different from a human's.

2.4.1 RL in Verifiable Domains (e.g. Math & Code)

In domains with a clear right or wrong answer (like math or coding), the process of RL is straightforward. For a single prompt (e.g., a math problem), the model generates many different potential solutions (e.g., 100 different chains of thought). A model then automatically checks which of these solutions arrive at the correct final answer, and the final step is to reinforce, i.e. the model is then trained and its parameters are updated to increase the probability of generating the successful token sequences. A key paper in this area, DeepSeek-R1 (DeepSeek-AI et al., 2025), showed that this process leads models to distribute their "thinking" over many tokens: It writes out intermediate steps, re-evaluates its work (e.g. "Wait, let me double-check that"), and tries multiple approaches. So, RL discovers reasoning strategies that are effective for models' architectures, even if a human wouldn't have written them that way.

2.4.2 RL in Unverifiable Domains like writing a joke (RLHF)

For creative tasks like writing a joke, where there's no single "correct" answer, Reinforcement Learning happens from Human Feedback (RLHF). The model generates several responses to a prompt (e.g., X different jokes). Then a human labeller ranks these responses from best to worst. However, instead of using humans to directly grade the LLM's performance billions of times for each sequence of responses, we use humans for a much smaller task to rank maybe 10,000 to 100,000 sets of responses and eventually train an AI judge model that can instantly rank given responses based on how a human would likely perceive it (by using same pre-training and post-training steps as mentioned above for LLM models). This judge, called a Reward Model, is a separate neural network and it can be used billions of times automatically. Its only job is to learn to predict human's preferences and is rewarded for getting the relative ordering correct.

Limitations of RLHF: While powerful, RLHF has a significant flaw. The reward model is just a simulation of a human, and like any neural network, it can be "gamed." If you run the RL process for too long, the LLM will find "adversarial examples"—nonsensical outputs that trick the reward model into giving a high score. For this reason, RLHF is more of a "fine-tuning" step to slightly improve model alignment and helpfulness, rather than a method for indefinite capability improvement like the RL used in verifiable domains. This is what the video means when it says "RHLF is not RL in the magical sense," unlike the systems like AlphaGo (Silver et al., 2016) which could practice indefinitely against a perfect game simulator to achieve superhuman performance.

11

2.5 Conclusion

The creation of a large language model is a multi-stage process, transforming a raw internet text predictor (base model) into a sophisticated, reasoning assistant.

- Pre-training gets training data and builds knowledge by predicting the next token across trillions of examples.
- SFT teaches the model how to converse like a helpful assistant by imitating human-written conversations.
- RL allows the model to move beyond imitation and discover effective, sometimes novel, reasoning strategies through trial-and-error practice.

However, it is important to remember the model's nature. It is a stochastic, token-based system with a finite computational budget for each step. Its knowledge has "mistakes," and its reasoning can be broken. It does not "know" things in the human sense but is a master of statistical pattern-matching. As such, LLMs should be treated as incredibly powerful tools—for brainstorming, for first drafts, for summarising, and for accelerating work—but their output should be verified by a human user who remains ultimately responsible for the final product.

References

Joshua Achiam. Vanilla policy gradient — spinning up documentation, 2018. URL https://spinningup.openai.com/en/latest/algorithms/vpg.html#pseudocode.

Kevin Afachao. Deep dive into reinforcement learning: Policy gradient algorithms. *Medium*, 2023. URL https://medium.com/@thekevin.afachao/deep-dive-into-reinforcement-learning-policy-gradient-algorithms-bbeabe2fc989. Accessed: 2025-05-23.

DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, et al. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning, 2025.

Hugging Face. Byte-Pair Encoding (BPE) Tokenization, in: The LLM Course. https://huggingface.co/learn/llm-course/en/chapter6/5, 2025.

Andrej Karpathy. Deep dive into llms like chatgpt, 2025. URL https://www.youtube.com/watch?v=zjkBMFhNj_g.

Sergey Levine. Policy gradients (uc berkeley cs 285 lecture notes), 2021. URL https://rail.eecs.berkeley.edu/deeprlcourse/.

MIT-6.7920. Policy gradient. https://web.mit.edu/6.7920/www/lectures/L13&14-2024fa-PolicyGradient.pdf, 2024. Lecture notes for MIT 6.7920: Reinforcement Learning, Fall 2024.

OpenAI. What are tokens and how to count them? https://help.openai.com/en/articles/4936856-what-are-tokens-and-how-to-count-them, 2025.

OpenAI, Achiam. Spinning up in deep rl. https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html, 2018. Accessed: 2025-05-23.

Guilherme Penedo, Hynek Kydlíček, Loubna Ben allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro Von Werra, and Thomas Wolf. The fineweb datasets: Decanting the web for the finest text data at scale. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024. URL https://openreview.net/forum?id=n6SCkn2QaG.

John Schulman. *Optimizing Expectations: From Deep Reinforcement Learning to Stochastic Computation Graphs*. PhD thesis, University of California, Berkeley, 2016. Specifically Chapter 2, as referenced in OpenAI Spinning Up documentation.

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4):229–256, 1992. doi: 10.1007/BF00992696. URL https://link.springer.com/article/10.1007/BF00992696.